

# High Performance Computing with GPU

---

## **Semester Project**

# Group Information – Section B

---

- **Group Members**

- Hamza Saleem 22i-1185
- Abdullah Mehmood 22i-1167
- Aneeq Ahmed Malik 22i-0978

- **Problem under Discussion**

- MNIST Dataset: 60,000 training, 10,000 test images (28x28 pixels)
- Neural Network: 784 (input) → 128 (hidden) → 10 (output).
- Challenge: CPU sequential processing limits performance
- Solution: Leverage GPU parallelism for acceleration
- Approach: Four versions (V1: CPU, V2: Naive CUDA, V3: Optimized CUDA, V4: cuBLAS)

- **Github:** <https://github.com/Aneeq-Ahmed-Malik/Neural-Network-Acceleration-on-GPUs>

# Initial Implementation – CPU Version

- Description: C-based, processes images sequentially
- Components: Matrix multiplications: W1 (128x784), W2 (10x128)
- Stochastic Gradient Descent for forward/backward passes

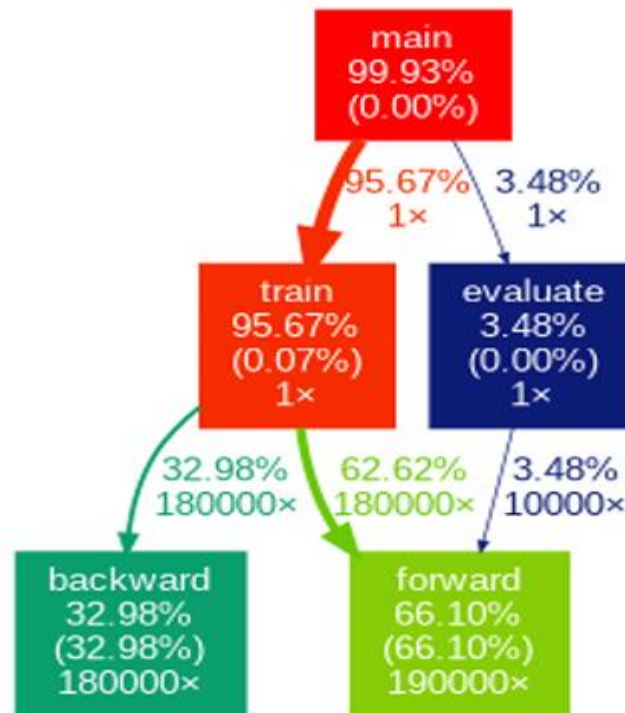


Figure 1: Call Graph of the Neural Network Application

# Towards V2 – Forward

// Forward pass

```
void forward{// params) {
```

```
    for (int i = 0; i < HIDDEN_SIZE; i++) {  
        hidden[i] = net->b1[i];  
        for (int j = 0; j < INPUT_SIZE; j++)  
            hidden[i] += net->W1[i][j] * input[j];  
    }
```

```
    relu(hidden, HIDDEN_SIZE);
```

```
    for (int i = 0; i < OUTPUT_SIZE; i++) {  
        output[i] = net->b2[i];  
        for (int j = 0; j < HIDDEN_SIZE; j++)  
            output[i] += net->W2[i][j] * hidden[j];  
    }
```

```
    softmax(output, OUTPUT_SIZE);
```

```
}
```

compute\_hidden

compute\_output

softmax

# Towards V2 – Backward

```
void backward(//params) {
```

```
for (int i = 0; i < OUTPUT_SIZE; i++)  
    d_output[i] = output[i] - target[i];
```

d\_hidden

```
for (int i = 0; i < HIDDEN_SIZE; i++) {  
    d_hidden[i] = 0;  
    for (int j = 0; j < OUTPUT_SIZE; j++)  
        d_hidden[i] += net->W2[j][i] * d_output[j];  
    d_hidden[i] *= (hidden[i] > 0);  
}
```

d\_output

```
for (int i = 0; i < OUTPUT_SIZE; i++)  
    for (int j = 0; j < HIDDEN_SIZE; j++)  
        net->W2[i][j] -= LEARNING_RATE * d_output[i] * hidden[j];  
for (int i = 0; i < OUTPUT_SIZE; i++)  
    net->b2[i] -= LEARNING_RATE * d_output[i];
```

out  
weights


```
for (int i = 0; i < HIDDEN_SIZE; i++)  
    for (int j = 0; j < INPUT_SIZE; j++)  
        net->W1[i][j] -= LEARNING_RATE * d_hidden[i] * input[j];  
for (int i = 0; i < HIDDEN_SIZE; i++)  
    net->b1[i] -= LEARNING_RATE * d_hidden[i];
```

hid  
weights

```
}
```

# Version – 2: Performance

Performance Comparison: V1 to V2




*Values show speedup factor compared to V1 baseline*

# Towards V3.1: Shared Memory (1)

```
__global__ void compute_hidden_layer(/* params */) {  
    int i = threadIdx.x;  
    if (i < HIDDEN_SIZE) {  
        double sum = net->b1[i];  
        for (int j = 0; j < INPUT_SIZE; j++)  
            sum += net->W1[i * INPUT_SIZE + j] * input[j];  
        hidden[i] = (sum > 0.0) ? sum : 0.0;  
    }  
}
```

Note: 2 more kernels follow the same pattern

```
__global__ void compute_output_layer(/* params */) {  
    int i = threadIdx.x;  
    if (i < OUTPUT_SIZE) {  
        double sum = net->b2[i];  
        for (int j = 0; j < HIDDEN_SIZE; j++)  
            sum += net->W2[i * HIDDEN_SIZE + j] * hidden[j];  
        output[i] = exp(sum);  
    }  
}
```



# Towards V3.1: Shared Memory (2)

```
__global__ void compute_hidden_layer(//params) {
```

```
int i = threadIdx.x;
```

Local Id

```
__shared__ double s_input[INPUT_SIZE];  
s_input[i] = input[i];
```

SharedMem

```
__syncthreads();
```

Local sync

```
if (i < HIDDEN_SIZE) {  
    double sum = net->b1[i];  
    for (int j = 0; j < INPUT_SIZE; j++)  
        sum += net->W1[i * INPUT_SIZE + j] *  
  
    hidden[i] = (sum > 0.0) ? sum : 0.0;  
}
```

s\_input[j];


Use Shared

```
}  
Note: Other kernels  
follow the same pattern
```



# Version – 3.1: Performance

Performance Comparison: V1 to V3.1



*Values show speedup factor compared to V1 baseline*

# Towards V3.2: Cuda Streams (1)

```
void train(// params) {  
    double* hidden, *output  
    double* d_hidden, *d_output, *d_input, *d_label;  
    // cuda Malloc & malloc for all device/host variables
```

```
    for (int epoch = 0; epoch < EPOCHS; epoch++) {  
        for (int i = 0; i < numImages; i++) {
```

```
            cudaMemcpy(d_input, images[i], size_t, H2D);  
            forward_cuda(net, d_input, d_output, d_hidden);  
            cudaMemcpy(d_label, labels[i], size_t, H2D);
```

dependent

```
            backward_cuda(net, d_input, d_hidden, d_output, d_label);  
            cudaMemcpy(output, d_output, size_t, D2H);
```

```
            // Compute loss & accuracy
```

```
        }
```

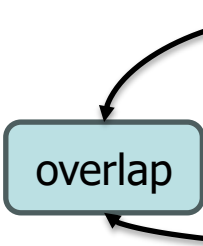
```
    }
```

```
}
```

But d\_input & d\_label of  
next iter can overlap  
backward of current iter

# Towards V3.2: Cuda Streams (2)

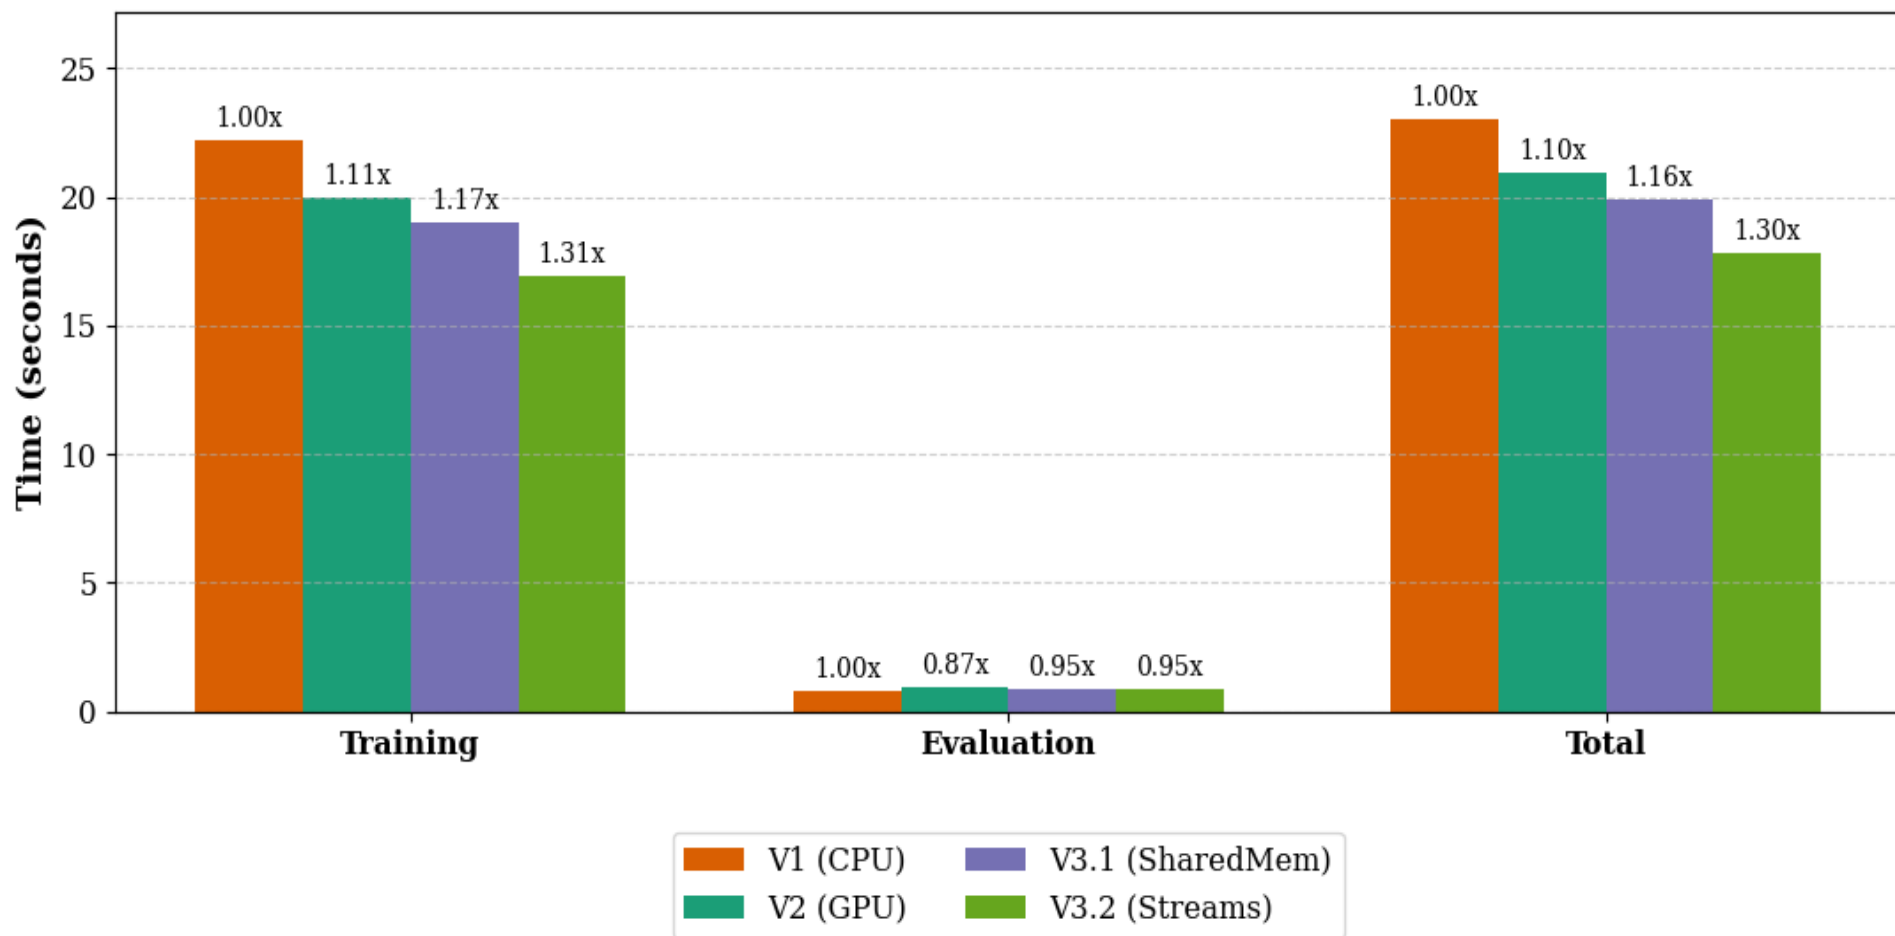
```
void train(/* params */) {  
    cudaMemcpyAsync(d_input, images[0], size, H2D, streams[0]);  
    cudaMemcpyAsync(d_label, labels[0], size, H2D, streams[0]);  
  
    for ( /* epoch loop */ ) {  
        forward_cuda(/*same*/, streams[0]);  
        for ( /* numImages loop */ ) {  
            int curr_str = i % 2, next_str = (i + 1) % 2;  
  
            backward_cuda(/*same*/, streams[curr_str]);  
            cudaMemcpyAsync(/* output */, streams[curr_str]);  
  
            if (i + 1 < numImages) {  
                cudaMemcpyAsync(/* input */, streams[next_str]);  
                cudaMemcpyAsync(/* label */, streams[next_str]);  
            }  
  
            cudaStreamSynchronize(streams[curr_str]);  
            forward_cuda(/*same*/, streams[next_str]);  
        }  
    }  
}
```



overlap

# Version – 3.2: Performance

Performance Comparison: V1 to V3.2



Values show speedup factor compared to V1 baseline

## Towards V3.3: Batch Processing (2)

---

- Problem: Still processing one image at a time → high overhead.
- Idea : Process multiple images simultaneously in a batch (e.g., 64 images).
- Implementation:
  - Allocate memory for entire batch.
  - Launch forward and backward CUDA kernels for the batch.
  - Use CUDA streams to overlap batches.
- Impact:
  - Reduces kernel launch overhead.
  - Increases GPU utilization.
  - Significant speedup compared to single image processing.
- Drawback:
  - Small loss in accuracy at large batch sizes.
  - Less frequent weight updates (compared to per-image SGD).

# Towards V3.3: Batch Processing (2)

```
void train(params) {
```

```
    for (int b = 0; b < min(BATCH_SIZE, numImages); b++) {  
        cudaMemcpyAsync(d_input[0] + b * INPUT_SIZE, /*same*/);  
        cudaMemcpyAsync(d_label[0] + b * NUM_CLASSES, /*same*/);  
    }
```

```
    for (/*epoch loop*/) {  
        forward_cuda(/*same*/, streams[0]);  
        for (int i = 0; i < numImages; i += BATCH_SIZE) {  
            int batch_size = min(BATCH_SIZE, numImages - i);
```

Batch  
Copy

```
            if (i + BATCH_SIZE < numImages) {  
                int next_batch_size = min(BATCH_SIZE,  
                    numImages - (i + BATCH_SIZE));
```

Batch  
Copy

```
                for (int b = 0; b < next_batch_size; b++) {  
                    cudaMemcpyAsync(/* input */, streams[next_str]);  
                    cudaMemcpyAsync(/* label */, streams[next_str]);  
                }
```


```
            }
```

```
        }
```

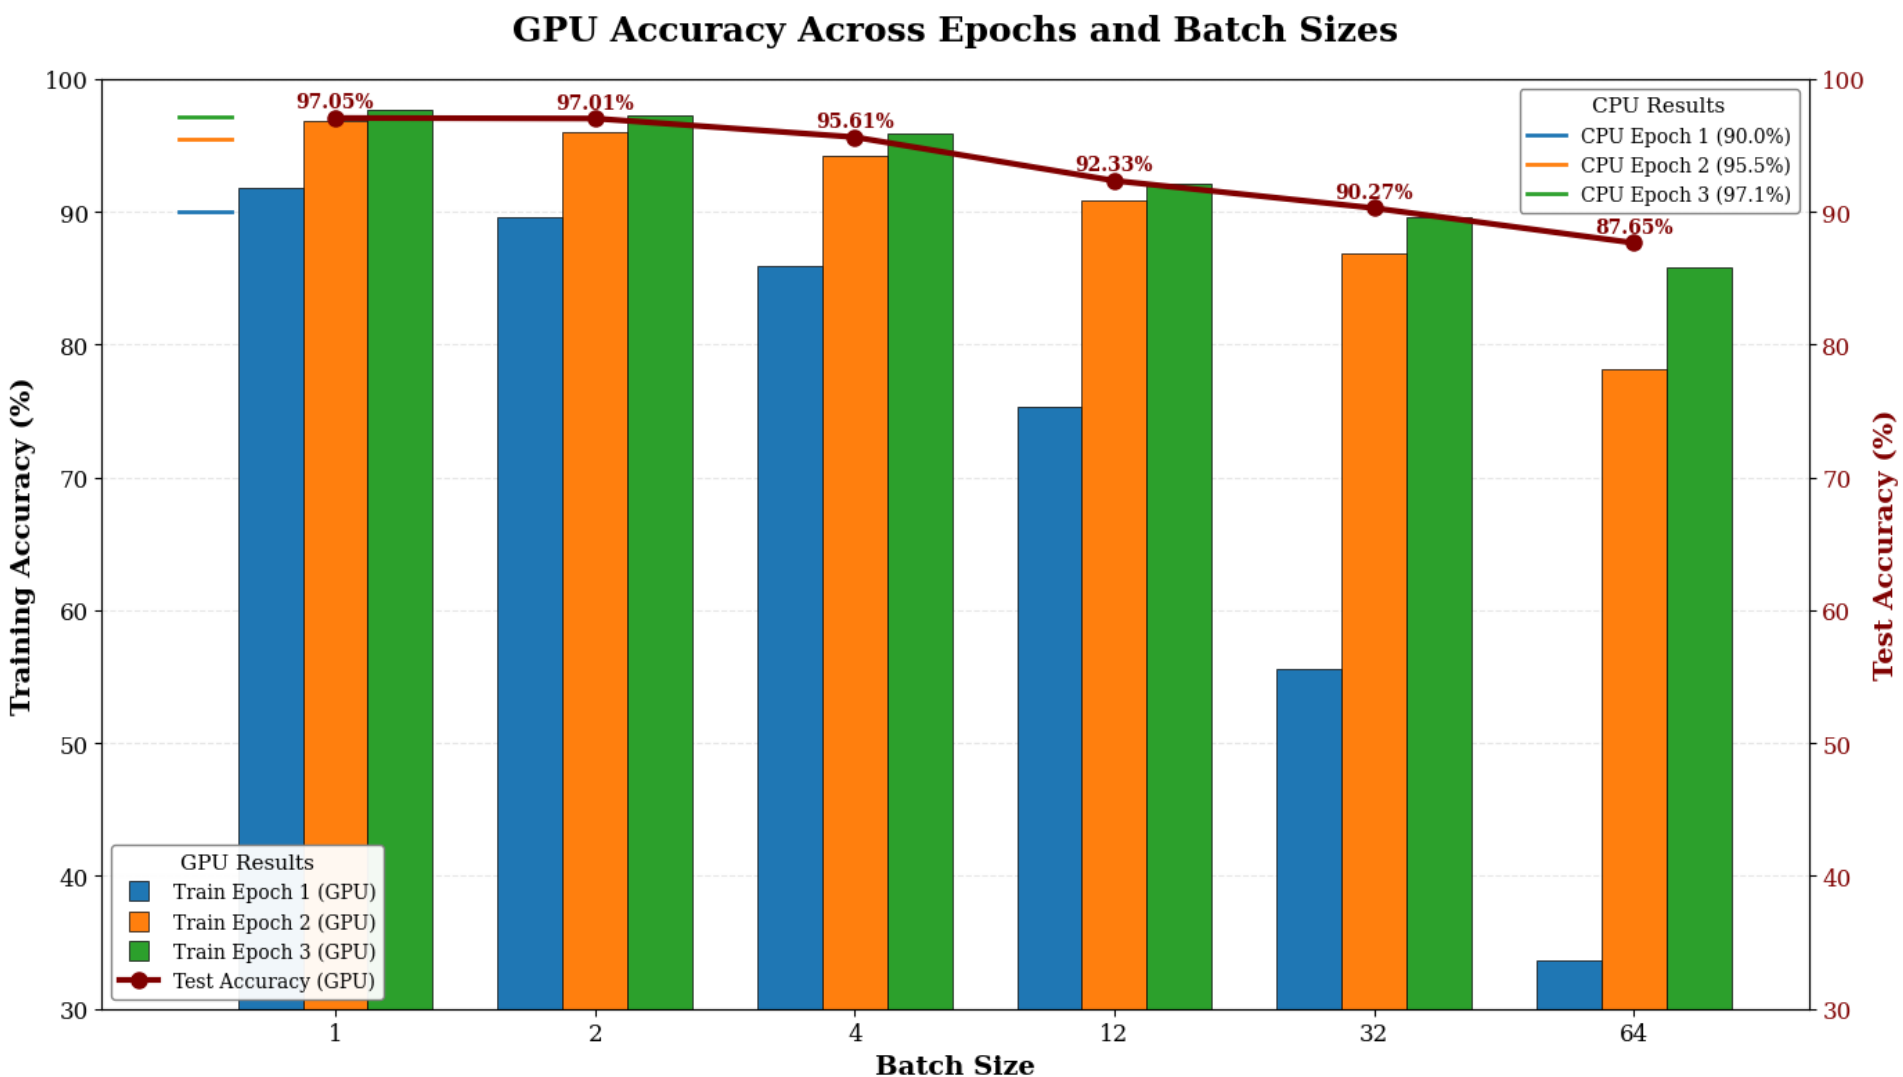
```
    }
```

```
}
```

# Version – 3.3: Performance



# Version – 3.3: Drawback



Note: CPU accuracies shown as small horizontal lines near y-axis.




# Towards V3.4: FP32 Processing

---

- Problem: Double precision (double) variables consume more memory and cause slower memory transfers on GPU.
- Idea: Switch to single precision (float, FP32) for faster computation and lower memory usage.
- Implementation:
  - Templatized the code: `(template<typename T>)`, so they work for both float and double.
- Impact:
  - Reduced memory consumption.
  - Faster memory transfers between CPU and GPU.
  - Increased kernel execution speed.
  - Improved training time significantly.
- Drawback:
  - Minor risk of numerical precision loss (but no noticeable accuracy drop for MNIST).

# Version – 3.4: Performance

Performance Comparison: V1 to V3.3



Values show speedup factor compared to V1 baseline

# Towards V3.5: Evaluation Kernel (1)

```
void evaluate(// params) {
    int correct = 0;
    double* hidden = (double*)malloc(sizeof(double) * HIDDEN_SIZE);
    double* output = (double*)malloc(sizeof(double) * OUTPUT_SIZE);
    double* d_hidden, *d_output, *d_input;
    // initialize device memory

    for (int i = 0; i < numImages; i++) {

        cudaMemcpy(d_input, images[i], size, H2D);
        forward_cuda(net, d_input, d_output, d_hidden, 0);
        cudaMemcpy(output, d_output, size, D2H);

        int pred = 0, actual = 0;
        for (int j = 0; j < output_size; j++) {
            if (output[j] > output[pred]) pred = j;
            if (labels[i][j] > labels[i][actual]) actual = j;
        }
        if (pred == actual) correct++;
    }
}
```

No  
dependency

## Towards V3.5: Evaluation Kernel (2)

```
void evaluate(// params) {
    int correct = 0;
    T* output = (T*)malloc(sizeof(T) * OUTPUT_SIZE * numImages);
    T *d_input, *d_hidden, *d_output;
    cudaMalloc((void**)&d_input, sizeof(T) * INPUT_SIZE * numImages);
    cudaMalloc((void**)&d_hidden, sizeof(T) * HIDDEN_SIZE * numImages);
    cudaMalloc((void**)&d_output, sizeof(T) * OUTPUT_SIZE * numImages);

    for (int i = 0; i < numImages; i++)
        cudaMemcpyAsync(d_input + i * INPUT_SIZE, images[i], size, H2D);

    forward_cuda_eval(net, d_input, d_output, d_hidden, numImages);

    // Copy full output back to host
    cudaMemcpy(output, d_output, size , D2H);


    // Accuracy calculation
    for (int i = 0; i < numImages; i++) {
        // same check pred
    }
}
```

```
graph TD; A(numImages) --> B(BatchSize)
```

BatchSize

# Version – 3.5: Performance

Execution Time Comparison: V1 to V3.4

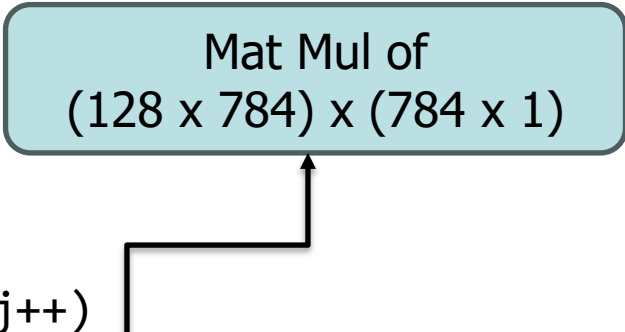


Values above bars show speedup factor compared to V1 baseline

# Towards V4: Tensor Cores

- Problem: Even after optimization, custom CUDA kernels were slower than what optimized libraries could achieve.
- Idea: Replace manual matrix multiplications with highly-optimized cuBLAS library functions.

```
__global__ void compute_hidden_layer(params) {  
    int i = threadIdx.x;  
    __shared__ double s_input[INPUT_SIZE];  
    s_input[i] = input[i];  
    __syncthreads();  
  
    if (i < HIDDEN_SIZE) {  
        double sum = net->b1[i];  
        for (int j = 0; j < INPUT_SIZE; j++)  
            sum += net->W1[i * INPUT_SIZE + j] * s_input[j];  
        hidden[i] = (sum > 0.0) ? sum : 0.0;  
    }  
}
```



Mat Mul of  
(128 x 784) x (784 x 1)

# Towards V4: Tensor Cores

- Implementation:
  - Used `cublasSgemm()` to perform fast batched matrix multiplications
  - It computes  $C = \alpha \times (A \times B) + \beta \times C$  (column major)
  - A and B are input matrices, C is the output matrix
  - alpha and beta are scaling factors (usually 1 and 0).
  - Here,  $C_{128 \times 1} = A_{128 \times 784} \times B_{784 \times 1}$

```
const T alpha = 1.0f;  
const T beta = 0.0f;  
cublasSgemm(handle,
```

```
    CUBLAS_OP_T, CUBLAS_OP_T,  
    HIDDEN_SIZE, BATCH_SIZE, INPUT_SIZE,  
    &alpha,  
    h_net.W1, INPUT_SIZE,  
    d_input, BATCH_SIZE,  
    &beta,  
    d_hidden, HIDDEN_SIZE);
```


Transpose input

Input Mat (A,B)

Output Mat (C)

# Version – 4: Performance

Performance Optimization: CPU to Tensor Cores



Numbers above bars show speedup relative to V1 (CPU version)



# Any Question So Far?

---

